

Q. Write a python program to import and export data using pandas library function

Import pandas as pd

```
airline_data = pd.read_csv("data/listings-austin.csv")
airline_data.head()
```

Output :-

Imported the pandas library into our environment
Printed the first 5 rows of data frame

* Reading the data from URL

```
URL = "https://archive.ics.uci.edu/ml/machine-learning-
      -databases/iris/iris.data"
```

```
col_names = ["Sepal-length-in-cm",
             "Sepal-width-in-cm",
             "Petal-length-in-cm",
             "Petal-width-in-cm",
             "Class"]
```

```
iris_data = pd.read_csv(URL, names=col_names)
iris_data.head()
```

Output :-

| | Sepal-length-in-cm | Sepal-width-in-cm | Petal-length-in-cm | Petal-width-in-cm |
|---|--------------------|-------------------|--------------------|-------------------|
| 0 | 5.1 | 3.5 | 1.4 | 0.2 |
| 1 | 4.9 | 3.0 | 1.6 | 0.2 |
| 2 | 4.7 | 3.2 | 1.3 | 0.2 |
| 3 | 4.6 | 3.1 | 1.5 | 0.1 |
| 4 | 5 | 3.6 | 1.4 | 0.2 |

→ exporting dataframe to CSV file

it is -> `df.to_csv('cleaned_iris_data.csv')`

Kaggle

→ Creating data frame

Pd. Dataframe ([{'Yes': [50, 21], 'No': [131, 213]}])

Output:

| | Yes | No |
|---|-----|-----|
| 0 | 50 | 131 |
| 1 | 21 | 2 |

→ Series

Pd. Series ([1, 2, 3, 4, 5])

| 0 | 1 |
|---|---|
| 1 | 2 |
| 2 | 3 |
| 3 | 4 |
| 4 | 5 |

dtype : int64

→ Reading Data Files

data = pd.read_csv("college-data.csv")

data = Pd. read_csv("college-data.csv")

data.head()

Returns first 5 data value of data set

* end to end project Hands on machine learning

Step 1:

get the data

data.head()

- gives first five data entries

data.info()

Output:

| | col / attr | non - null |
|---|-------------------------|------------|
| 0 | longitude | 20640 |
| 1 | latitude | 20640 |
| 2 | housing - age | 20640 |
| 3 | total - rooms | 20640 |
| 4 | total - bedrooms | 20640 |
| 5 | population | 20640 |
| 6 | households | 20640 |
| 7 | median - income | 20640 |
| 8 | median - house - income | 20640 |
| 9 | ocean - proximity | 20640 |

data['ocean - proximity'].value_counts()

data.describe()

data.hist(bins = 50, figsize = (10, 15))

plt.show()

Step 2:

Discover & visualize the data to gain insights.

```
data.plot(kind = 'scatter', x = 'longitude', y = 'latitude')
```

plt.show()

```
data.plot(kind = 'scatter', x = 'longitude', y = 'latitude',
          alpha = 0.1)
```

plt.show

```
data.plot(kind = 'scatter', x = 'longitude', y = 'latitude',
          alpha = 0.1)
```

plt.show(2)

~~data[['population', 'median_house_value']].corr()~~
~~corr_matrix = housing.corr().abs()~~
~~corr_matrix['median_house_value'].sort_values(ascending = True)~~
~~data.plot(kind = 'scatter', x = 'median_income', y = 'median_house_value', figsize = (12, 8), alpha = 0.1)~~

plt.show(1)

Step 3: prepare the data for machine learning
dgb.

```
housing.drop(['ocean_proximity'], axis = 1)
```

housing_cat = housing[[' Ocean_proximity']]

housing_cat = pd.get_dummies(housing_cat)

imputer = Imputer(strategy='median')
imputer.fit(housing_cat)

train_index, test_index = train_test_split(housing_cat, test_size=0.2, random_state=42)
y = housing['income_cat']

train_set = housing_cat.loc[train_index]

test_set = housing_cat.loc[test_index]

housing = train_set.copy()

5. Select one train model

from sklearn.linear_model import LinearRegression

lin_reg = LinearRegression()

lin_reg.fit(X= housing_prepred, y=housing_labels)

some_data = housing.loc[:5]

print(lin_reg.predict(some_data))

lin_mse = mean_squared_error(housing_labels, lin_reg.predict(housing_prepred))

Simple linear regression

Import pandas as pd

Import numpy as np

Import matplotlib.pyplot as plt

Import seaborn as sns

from sklearn import linear_model import linearregression

df_sd = pd.read_csv('salary_data.csv')

df_sd.head()

plt.title('Salary Plot')

sns.distplot(df_sd['salary'], kde=True)

plt.show()

plt.scatter(df_sd['experience'], df_sd['salary'])

df_sd['salary'].columns = ('high')

plt.title('Salary vs experience')

plt.show()

Split data

x = df_sd['high'][:, :1]

y = df_sd['high'][:, 1:]

split into train & test sets

x_train, x_test, y_train, y_test = train_test_split

x, y, test_size=0.2, random_state=0)

train model

regressor = LinearRegression()

regressor.fit(x_train, y_train)

$y_pred_test = \text{regressor.predict}(X\text{-test})$
 $y_pred_train = \text{regressor.predict}(X\text{-train})$

visual predictions

plt.scatter(X-train, y-train, color='light red')
plt.plot(X-train, y-pred-train)

plt.show

coefficient and intercept

print(lm1.coef[0]) # coefficient = regressor.coef[0]
print(lm1.intercept) # regressor.intercept[0]

coefficient: [0.9372 - 875(2)]

intercept: [26780.099]

Multiple linear regression

Import pandas as pd

Import numpy as np

Import seaborn as sns

```
from sklearn.linear_model import LinearRegression
```

```
df_store = pd.read_csv('store.csv')
```

```
df_store = df_store.drop(['id'])
```

```
df.describe()
```

Distribution

```
plt.title('Profit distribution plot')
```

```
sns.distplot(df_store['profit'])
```

```
plt.show()
```

Relationship between profit and R&D spend

```
plt.scatter(df_store['R&D Spend'], df_store['profit'])
```

```
plt.show()
```

Split data

```
X = df_store[:, :-1].values
```

```
y = df_store[:, -1].values
```

Split into train data

x_train, x-test, y-train, y-test = train-test-split
(x, y, test_size = 0.2, random_state = 0)

Train model

regressor = LinearRegression()

regressor.fit(x-train, y-train)

Predict Results

y-pred = regressor.predict(x-test)

Decision tree using entropy

Import pandas as pd

iris_data = pd.read_csv('Iris.csv')

iris_data.head()

Import pandas as pd

from sklearn.model_selection import train_test_split

from sklearn.tree import DecisionTreeClassifier

from sklearn.metrics import accuracy_score

from sklearn.tree import DecisionTreeClassifier, plot_tree

import matplotlib.pyplot as plt

iris_data.drop(['Id'], inplace=True, axis=1)

X = iris_data.drop(['Species'], axis=1)

y = iris_data['Species']

X_train, X_test, y_train, y_test = train_test_split

(X, y), test_size=0.4, random_state=42)

dt_classifier = DecisionTreeClassifier(criterion='entropy',
random_state=42)

dt_classifier.fit(X_train, y_train)

y_pred = dt_classifier.predict(X_test)

accuracy = accuracy_score(y_test, y_pred)

print(f'accuracy: {accuracy}')

• polypropylene $\epsilon = 2.45$

entropy $= 1.581$

Sample = 98

Value = [27, 31, 32,]

clm = 1115 - 1115

entropy $= 0.0$

Sample = 27

Value $\approx [27, 0, 0]$

clm = 1115 - 1115

polywidth cm $\epsilon = 1.75$

entropy $= 1.6$

Sample = 83

Value $\approx [37, 32, 2]$

clm = 1115 - 1115

\$ Build logistic regression model for a given dataset

import pandas as pd

from matplotlib import pyplot as plt
%matplotlib inline

df = pd.read_csv("insurance.csv")
df.head()

plt.scatter(df['age'], df['bought_insurance'], marker='+',
 color='green', label='bought insurance')
plt.scatter(df['age'], df['bought_insurance'], marker='x',
 color='red', label='did not buy insurance')

from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split
(df[['age']], df['bought_insurance'],
train_size=0.8)

Model = LogisticRegression()
model.fit(X_train, y_train)

y_predicted = model.predict(X_test)

print(y_predicted)
[1 0 1 0]

print(Xtest)

| | Age |
|----|-----|
| 9 | 61 |
| 23 | 45 |
| 21 | 26 |

4 46
7 70
8 72

Print ("Model Accuracy"), model.score (X-test)

model accuracy 0.833

print ("coefficient", model.coef_)
print ("intercept", model.intercept_)

coefficient [[0.1499]]
intercept [-5.6031]

Import math

def sigmoid(z):
 return 1 / (1 + math.exp(-z))

def prediction_function(age):

$$z = 0.642 * \text{age} - 1.53$$

$$y = \text{Sigmoid}(z)$$

return y

$$\text{age} = 35$$

prediction_function(35)

$$0.485004$$

$$\text{age} = 43$$

prediction_function(43)

$$0.568565$$

* SVM & KNN

Import random as rd

from sklearn.model_selection import train - test

from sklearn.metrics import accuracy_score

from sklearn.tree import plot_tree

import matplotlib.pyplot as plt

iris_data = pd.read_csv('iris.csv')

iris_data.head()

label_mapping = { 'Iris-setosa': 1, 'Iris-versicolor': 2, 'Iris-virginica': 3 }

iris_data['Species'] = iris_data['Species'].map(label_mapping)

plt.Scatter(iris_data['Sepal length(cm)'], iris_data['Sepal width(cm)'], c=iris_data['Species'], cmap='viridis')

plt.xlabel('length (cm)')

plt.ylabel('width (cm)')

plt.title('Scatter plot of iris dataset')

plt.colorbar(label='Species')

plt.show()

iris_data.drop(['Id'], inplace=True, axis=1)

x = iris_data.drop(['Species'], axis=1)

y = iris_data['Species']

x_train, x_test, y_train, y_test = train - test split
(n, y, test_size=0.3, random_state=92)

$\text{SVM - Clasifier} = \text{SVC}$

$\text{SVM - Clasifier . fit}(X_{\text{train}}, y_{\text{train}})$

fitting with multi class will give

$y_{\text{predicted}} = \text{SVM - Clasifier . predict}(X_{\text{test}})$

accuracy = $\text{SVM - Clasifier . predict}(X_{\text{test}})$

accuracy = accuracy_score(y_{\text{test}}, y_{\text{pred}})

print("Accuracy", accuracy)

Accuracy : 0.9333333333333333

Accuracy : 0.9333333333333333

for KNN

accuracy : 0.9833

(accuracy, recall = accuracy_score(y_{\text{test}}, y_{\text{pred}}))

(accuracy, precision = precision_score(y_{\text{test}}, y_{\text{pred}}))

(accuracy, f1 = f1_score(y_{\text{test}}, y_{\text{pred}}))

(accuracy, support = support)

accuracy, recall, f1 = confusion_matrix(y_{\text{test}}, y_{\text{pred}})

accuracy, precision, f1 = precision_recall_f1_score(y_{\text{test}}, y_{\text{pred}})

accuracy, recall, f1 = classification_report(y_{\text{test}}, y_{\text{pred}})

accuracy, recall, f1 = confusion_matrix(y_{\text{test}}, y_{\text{pred}})

accuracy, precision, f1 = precision_recall_f1_score(y_{\text{test}}, y_{\text{pred}})

accuracy, recall, f1 = classification_report(y_{\text{test}}, y_{\text{pred}})

accuracy, recall, f1 = precision_recall_f1_score(y_{\text{test}}, y_{\text{pred}})

accuracy, recall, f1 = classification_report(y_{\text{test}}, y_{\text{pred}})

8. AdaBoost

```

from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import AdaBoostClassifier
from sklearn import metrics
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_

```

```
iris_data = pd.read_csv('iris.csv')
```

```
iris_data.head()
```

```
iris_data.info()
```

```
iris_data.drop(['Id'], inplace=True, axis=1)
```

```
X = iris_data.drop(['Species'], axis=1)
```

```
y = iris_data['Species']
```

```
iris_data.info()
```

```
X_train, X_test, y_train, y_test = train_test_
```

```
(X, y, test_size=0.3, random_state=42)
```

`MyLogRegModel` = logistic regression

`adaboost` = AdaBoostClassifier(n_estimators=150, learning_rate=0.01)

= `MyLogRegModel`, learning_rate=1)

```
model = adaboost.fit(X_train, y_train)
```

```
y_pred = model.predict(X_test)
```

```
metrics.accuracy_score(y_test, y_pred)
```

0.9633

* Random forest

Pip install scikit-learn

Import pandas as pd

```
from sklearn.model_selection import train_test_split  
from sklearn.ensemble import RandomForestClassifier  
from sklearn.datasets import load_iris  
from sklearn.metrics import accuracy_score
```

Iris-data = pd.read_csv('Iris.csv')

~~Iris-data~~

Iris-data.head()

Iris-data.info()

iris-data.drop(['Id'], inplace=True, axis=1)

X = iris-data.drop(['Species'], axis=1)

y = iris-data['Species']

X-train, X-test, y-train, y-test = train_test_split(
(X, y, test_size=0.4, random_state=42))

rf-classifier = RandomForestClassifier(n_estimators=100,
random_state=42)

rf_classifier = RandomForestClassifier(n_estimators=100,
n_estimators=100)

23/05/2023

23/05/2023

ANN

Import Numpy as np

$x = np.array([[[2, 4], [1, 3], [3, 6]]], dtype= float)$

$y = np.array([192, 86, 187]), dtype= float$

$x = x / np.max(x, axis= 0)$

$y = y / 100$

epoch = 5000

lr = 0.1

input layer - neurons = 2

hidden layer - neurons = 3

output - neurons = 1

$wh = np.random.uniform(size=(\text{input layer - neurons}, \text{hidden layer - neurons}))$

$b_h = np.random.uniform(size=(1, \text{hidden layer - neurons}))$

$wout = np.random.uniform(size=(\text{hidden layer - neurons}, 1))$

$b_o = np.random.uniform(size=(1, \text{output - neurons}))$

def sigmoid(z):

return 1 / (1 + np.exp(-z))

def derivative_sigmoid(z):

return z * (1 - z)

for i in range(epoch):

h_ip1 = np.dot(x, wh)

h_ip = h_ip1 + b_h

h_layer_act = sigmoid(h_ip)

out_ip1 = np.dot(h_layer_act, wout)

$$\text{out_link} = \text{out_link}_1 + \text{bias}$$

$$\text{output} = \text{Sigmoid}(\text{out_link})$$

$$E_0 = y - \text{output}$$

$$\text{outgrad} = \text{derivatives} - \text{Sigmoid}(\text{out_link})$$

$$d_{\text{output}} = E_0 * \text{outgrad}$$

$$\text{Middlegrad} = \text{derivatives} - \text{Sigmoid}(\text{ChLayer} - \text{act})$$

$$d_{\text{hiddenlayer}} = G_H * \text{middlegrad}$$

$$w_{0t} = \text{hiddenlayer} - T \cdot \text{dot}(d_{\text{output}}) + b_0$$

$$w_{ht} = Y - T \cdot \text{dot}(d_{\text{hiddenlayer}}) + b_h$$

```
Print ("Input : " + str(x))
```

```
Print ("Actual output : " + str(y))
```

```
Print ("Predicted output : " + str(output))
```

30 → M
 30 → M

K-means algorithm

Import matplotlib.pyplot as plt

from sklearn import datasets

from sklearn.cluster import KMeans

import pandas as pd

import numpy as np

iris = datasets.load_iris()

X = pd.DataFrame(iris.data)

X.columns = ['Sepal-length', 'Sepal-width',

'Petal-length', 'Petal-width']

y = pd.DataFrame(iris.target)

y.columns = ['Targets']

model = KMeans(n_clusters=3)

model.fit(X)

plt.figure(figsize=(4, 14))

colormap = np.array(['red', 'lime', 'blue'])

plt.subplot(2, 2, 1)

plt.scatter(x.Petal-length, x.Petal-width, c=colormap[y.Targets], s=40)

plt.title('Red Clinton')

plt.title('Petal length')

plt.title('Petal width')

plt.subplot(2, 2, 2)

plt.scatter(x.Petal-length, x.Petal-width, c=colormap)

[model - 1d k-], $s = 40$)

PLT - title ('K-means cluster')

PLT - xtitle ('Data length')

PLT - ytitle ('Pct width')

and then start with some code, and then
choose some file to open

(readfile)

(f'9023013001') in

and then make data and make plot

and then make plot

and then choose some file, and then

(readfile) in

(f'9023013002') in

(make a new, clean dataset)

(readfile) in

(make a new, clean dataset)

* PCA

Import matplotlib.pyplot as plt

Import pandas as pd

Import numpy as np

Import Seaborn as sns

```
from sklearn.datasets import load_breast_cancer  
cancer = load_breast_cancer()
```

cancer.data.shape()

Print (cancer['DESCR'])

```
df = pd.DataFrame(cancer['data'], columns=cancer['
```

df.head()

```
from sklearn.preprocessing import StandardScaler  
scaler = StandardScaler()
```

Scaler.fit(df)

```
Scaled-data = Scaler.transform(df)
```

from sklearn.decomposition import PCA

```
PCA = PCA(n_components=2)
```

PCA.fit(Scaled-data),

```
X_PCA = PCA.transform(Scaled-data)
```

Scaled-data.shape

X_PCA.shape
(569, 2)

Import matplotlib.pyplot as plt

plt.figure(figsize=(8,6))

plt.Scatter(XPCA[:,0], X_PCA[:,1], c=cancer
['target'], cmap='Plasma')

plt.xlabel('first principle component')

plt.ylabel('second principle component')

plt.title('PCA of breast cancer dataset')

plt.colorbar(label='cancer type')

plt.show()